



+


# 2D Convolution Performance Analysis

—

# The Algorithm

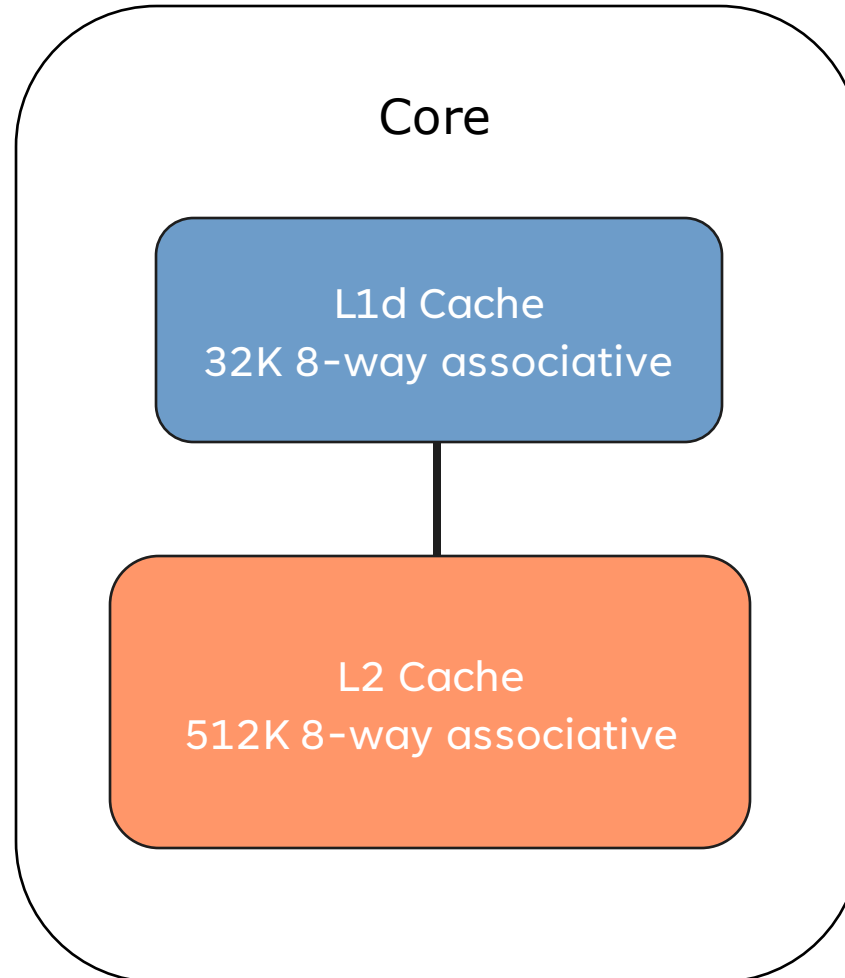
- 2D Convolution is an highly used algorithm in image filtering and Convolutional Neural Networks.
- Highly Parallelizable.
- CPU bound.

Given that the algorithm is parallelizable, let's try to implement a SIMT version and [analyse the speedup obtained](#).



# System

- Architecture : x86\_64
- CPU: AMD Ryzen 5 3500U
- Thread(s) per core: 2
- Cores : 4
- Frequency: 2.2 Ghz



# Measurements



- We measured **Wall clock time** using `std::chrono`

## Performance indexes

- Cumulative Speedup
- Execution time



# 1st Implementation


```
for(auto x = start; x < end; x++)
{
    for(auto y = 0; y < out_size_y; y++)
    {
        for(auto kx = 0; kx < kernel_size; kx++)
        {
            for(auto ky = 0; ky < kernel_size; ky++)
            {
                output[x * out_size_y + y] += input[(x + kx) * in_size_x + y + ky] * kernel[kx * kernel_size + ky];
            }
        }
    }
}
```

- Thread's task: compute a constant number of output's elements
- **output[]** is shared among all threads




# 1st Implementation: execution time

Mean exec time vs Number of threads - 1st implementation - (kernel=3x3)



# 1st Implementation: cumulative speedup

Cumulative speedup vs Number of threads - 1st implementation - (kernel=3x3)



# 2nd Implementation

```
// This holds the convolution results.
int convolute = 0;

// Fill output matrix: rows and columns are i and j respectively.
for (auto x = start; x < end; x++)
{
    for (auto y = 0; y < out_size_y; y++)
    {
        // Kernel rows and columns are kx and ky respectively.
        for (auto kx = 0; kx < kernel_size; kx++)
        {
            for (auto ky = 0; ky < kernel_size; ky++)
            {
                // Convolute here.
                convolute += (input[(x + kx) * in_size_x + (y + ky)] * kernel[kx * kernel_size + ky]);
            }
            // Add result to output matrix.
            output[x * out_size_y + y] = convolute;

            // Needed before we move on to the next index.
            convolute = 0;
        }
    }
}
```


- **convolute** is private for each thread.






# 2nd Implementation: execution time

Mean exec time vs Number of threads - 2nd implementation - (kernel=3x3)



# 2nd Implementation: cumulative speedup

Cumulative speedup vs Number of threads - 2nd implementation - (kernel=3x3)



# 2nd Implementation: analysis

convolute(void\*): **1st** implementation performance

Event Type	Incl.	Self	Short	Formula
Instruction Fetch	7 994 955 698	7 994 955 698	Ir	
L1 Instr. Fetch Miss	7	7	I1mr	
LL Instr. Fetch Miss	7	7	ILmr	
Data Read Access	3 955 577 810	3 955 577 810	Dr	
L1 Data Read Miss	4 191 748	4 191 748	D1mr	
LL Data Read Miss	2 096 130	2 096 130	DLmr	
Data Write Access	217 894 986	217 894 986	Dw	
L1 Data Write Miss	4	4	D1mw	
LL Data Write Miss	4	4	DLmw	
L1 Miss Sum	4 191 759	4 191 759	L1m = I1mr + D1mr + D1mw	
Last-level Miss Sum	2 096 141	2 096 141	LLm = ILmr + DLmr + DLmw	

convolute(void\*): **2nd** implementation performance

Event Type	Incl.	Self	Short	Formula
Instruction Fetch	5 179 135 252	5 179 135 252	Ir	
L1 Instr. Fetch Miss	7	7	I1mr	
LL Instr. Fetch Miss	7	7	ILmr	
Data Read Access	2 832 601 798	2 832 601 798	Dr	
L1 Data Read Miss	3 144 200	3 144 200	D1mr	
LL Data Read Miss	1 049 092	1 049 092	DLmr	
Data Write Access	100 569 136	100 569 136	Dw	
L1 Data Write Miss	1 047 558	1 047 558	D1mw	
LL Data Write Miss	1 047 558	1 047 558	DLmw	
L1 Miss Sum	4 191 765	4 191 765	L1m = I1mr + D1mr + D1mw	
Last-level Miss Sum	2 096 657	2 096 657	LLm = ILmr + DLmr + DLmw	

- L1 Data Read Misses have decreased by 25%.
- L3 Data Read Misses have decreased by roughly 50%.
- L1 Miss sum and L3 Miss sum have not changed.



# 3rd Implementation

```
// This holds the convolution results.
int convolute = 0;

// Fill output matrix: rows and columns are i and j respectively.
for (auto x = start; x < end; x++)
{
    for (auto y = 0; y < out_size_y; y++)
    {
        // Kernel rows and columns are kx and ky respectively.
        for (auto kx = 0; kx < kernel_size; kx++)
        {
            convolute += (input[(x + kx) * in_size_x + (y + 0)] * kernel[kx * kernel_size + 0]);
            convolute += (input[(x + kx) * in_size_x + (y + 1)] * kernel[kx * kernel_size + 1]);
            convolute += (input[(x + kx) * in_size_x + (y + 2)] * kernel[kx * kernel_size + 2]);
        }
        // Add result to output matrix.
        output[x * out_size_y + y] = convolute;


        // Needed before we move on to the next index.
        convolute = 0;
    }
}
```

- **Loop unrolling**
- Variables classified as **registers**




# 3rd Implementation: execution time

Mean exec time vs Number of threads - 3rd implementation - (kernel=3x3)



# 3rd Implementation: cumulative speedup

Cumulative speedup vs Number of threads - 3rd implementation - (kernel=3x3)



# 3rd Implementation: analysis

convolute(void\*): **2nd** implementation performance

Event Type	Incl.	Self	Short	Formula
Instruction Fetch	5 179 135 252	5 179 135 252	Ir	
L1 Instr. Fetch Miss	7	7	I1mr	
LL Instr. Fetch Miss	7	7	ILmr	
Data Read Access	2 832 601 798	2 832 601 798	Dr	
L1 Data Read Miss	3 144 200	3 144 200	D1mr	
LL Data Read Miss	1 049 092	1 049 092	DLmr	
Data Write Access	100 569 136	100 569 136	Dw	
L1 Data Write Miss	1 047 558	1 047 558	D1mw	
LL Data Write Miss	1 047 558	1 047 558	DLmw	
L1 Miss Sum	4 191 765	4 191 765	L1m = I1mr + D1mr + D1mw	
Last-level Miss Sum	2 096 657	2 096 657	LLm = ILmr + DLmr + DLmw	

convolute(void\*): **3rd** implementation performance

Event Type	Incl.	Self	Short	Formula
Instruction Fetch	2 028 089 892	2 028 089 892	Ir	
L1 Instr. Fetch Miss	5	5	I1mr	
LL Instr. Fetch Miss	5	5	ILmr	
Data Read Access	335 224 944	335 224 944	Dr	
L1 Data Read Miss	3 144 204	3 144 204	D1mr	
LL Data Read Miss	1 049 096	1 049 096	DLmr	
Data Write Access	16 760 856	16 760 856	Dw	
L1 Data Write Miss	1 047 558	1 047 558	D1mw	
LL Data Write Miss	1 047 558	1 047 558	DLmw	
L1 Miss Sum	4 191 767	4 191 767	L1m = I1mr + D1mr + D1mw	
Last-level Miss Sum	2 096 659	2 096 659	LLm = ILmr + DLmr + DLmw	

- Data Read Access have decreased by 89%.
- L1 Miss sum and L3 Miss sum have not changed.




# GPU Implementation

Each thread computes one element.

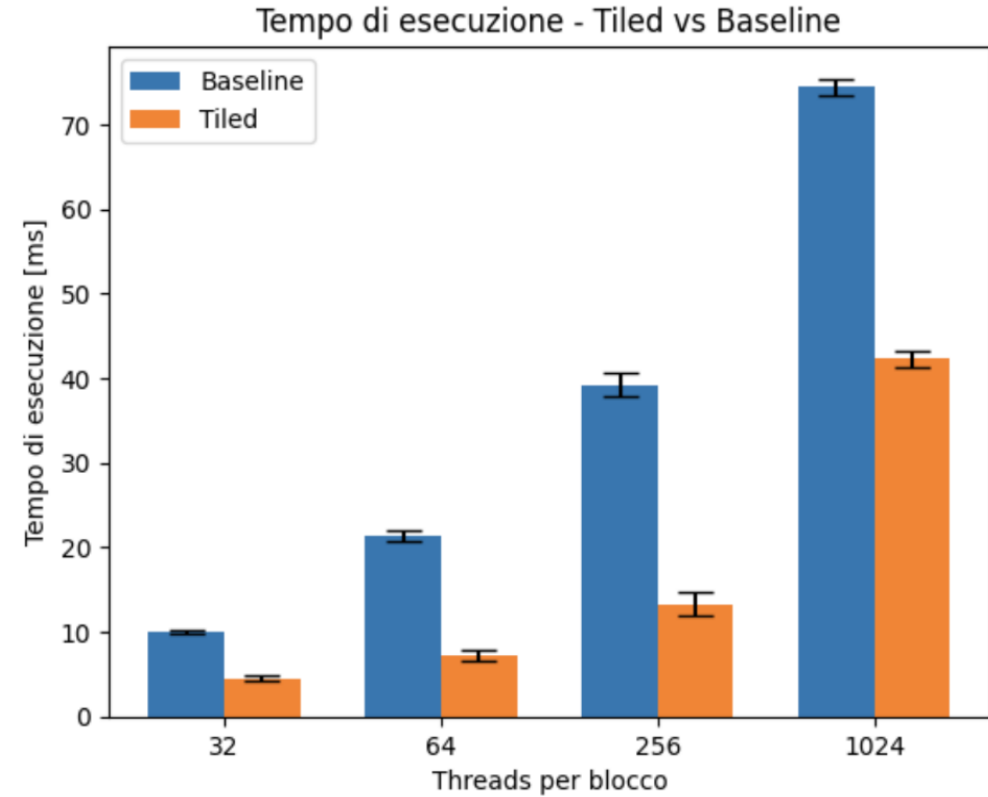
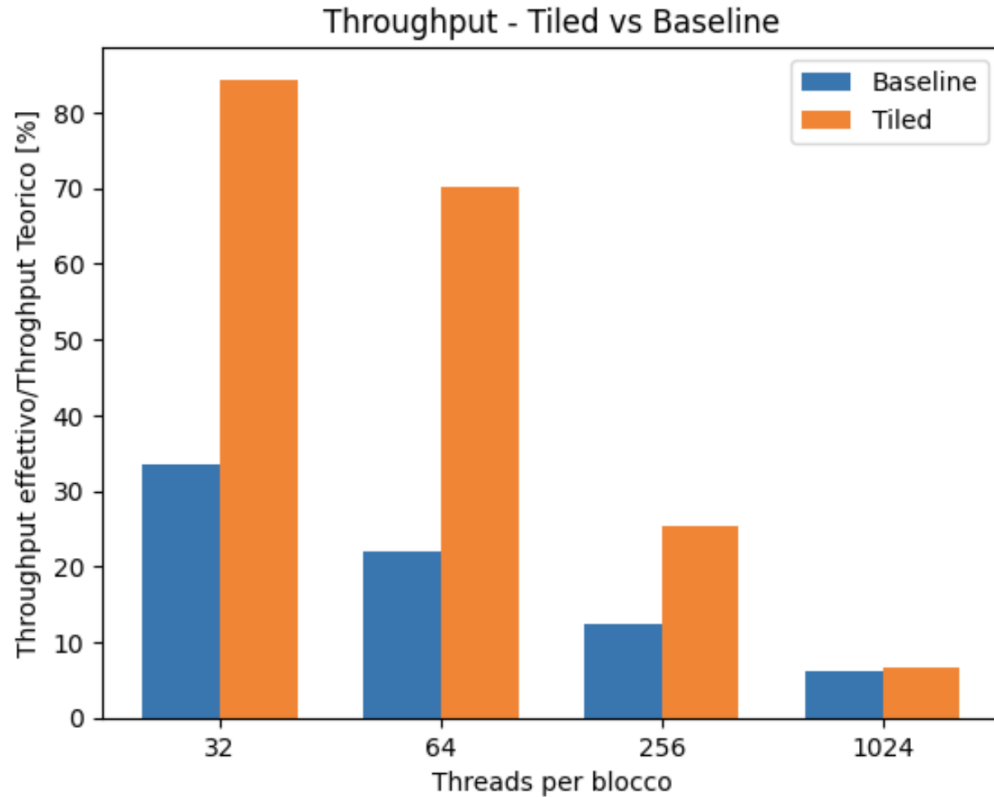
Two versions:

- Baseline
- Tiled (exploits shared memory)





# Performace comparison




## Observations:

- Shared version always performs better.
- Increasing block size worsen performance.



# Why Shared performs better

Base:



Shared:




We get **much less stalls on global memory loads** instruction




# Why increasing threads in a block worsen performance

– 64 Threads per block:



– 256 Threads per block:



- Number of stalled warp cycles doubles, that is in response to **contention on global memory**.



# Conclusions

---

- **CPU**

- If the **input size is small**, the advantages of **multithreading** are **nullified by** the **overhead** resulting from thread scheduling.
- Massive memory access, such as in convolution, may not allow for achieving speedups that are proportional to the number of threads.

- **GPU**

- The **performance improves** as the **block dimension decreases** when there is no need to exchange information between threads, as in our algorithm.
- Exploiting shared memory maximizes throughput.

